

Introduction to Pandas

- What is Pandas?

Pandas is a Python library widely used for data manipulation and analysis. It offers intuitive data structures like Series and DataFrame, akin to labeled arrays and tables, respectively. With Pandas, users can effortlessly perform tasks like cleaning, reshaping, and aggregating data. It seamlessly integrates with other Python libraries such as NumPy, Matplotlib, and Scikit-learn, making it a versatile tool for data analysis and visualization.

- Why Pandas for Data Science?

Pandas is widely preferred for data science due to several reasons:

Easy Data Manipulation: Pandas simplifies data manipulation tasks such as cleaning, filtering, transforming, and analyzing data.

Flexible Data Structures: With Series and DataFrame, Pandas provides flexible and intuitive data structures that can handle a wide variety of data types and formats, including structured, semi-structured, and time series data.

Integration with Other Libraries: Pandas seamlessly integrates with other Python libraries such as NumPy, Matplotlib, and Scikit-learn, enabling users to leverage the full ecosystem of Python tools for data analysis, visualization, and machine learning.

Rich Functionality: Pandas offers a rich set of functions and methods for data manipulation, including powerful tools for handling missing data, reshaping datasets, grouping and aggregating data, and performing time series analysis.

Community Support: Being an open-source project, Pandas benefits from a large and active community of users and contributors, ensuring continuous development, support, and improvement of the library.

- Installing Pandas

Installing Pandas is straightforward using Python's package manager, pip.

Open your terminal or command prompt and run the following command (you can copy & paste):

```
In [1]: # pip install pandas
```

This will download and install the latest version of Pandas along with any required dependencies.

Once installed, you can import Pandas into your Python scripts or notebooks using the following import statement:

```
In [2]: import pandas as pd
```

- Loading Data to Pandas

Pandas provides convenient functions to read data from various file formats into DataFrame objects. For this tutorial, we will focus on reading data from CSV files. However, Pandas also supports reading data from other formats such as Excel, JSON, SQL databases, and more.

```
In [3]: # Reading data from a CSV file into a DataFrame
df = pd.read_csv('clothing_prices.csv')
```

Data Exploration and Manipulation

Data Viewing and Accessing: head(), tail(), loc(), iloc()

- Display the first rows of the dataframe

```
In [4]: # By default, it returns the first 5 rows
df.head()
```

```
Out[4]:
```

	Brand	Category	Color	Size	Material	Price
0	New Balance	Dress	White	XS	Nylon	182.0
1	New Balance	Jeans	Black	XS	Silk	57.0
2	Under Armour	Dress	Red	M	Wool	127.0
3	NaN	Shoes	Green	M	Cotton	77.0
4	NaN	Sweater	White	M	Nylon	113.0

- Display the last rows of the dataframe

```
In [5]: # By default, it returns the last 5 rows
df.tail()
```

```
Out[5]:
```

	Brand	Category	Color	Size	Material	Price
995	Puma	Jeans	Black	NaN	Polyester	176.0
996	Puma	Jacket	Red	XXL	Silk	110.0
997	Reebok	Sweater	Blue	NaN	Denim	127.0
998	NaN	Sweater	Black	XXL	Denim	NaN
999	NaN	NaN	Yellow	XS	Wool	NaN

- Display a specific number of rows

```
In [6]: # e.g. return the first 4 rows
df.head(4)
```

```
Out[6]:
```

	Brand	Category	Color	Size	Material	Price
0	New Balance	Dress	White	XS	Nylon	182.0
1	New Balance	Jeans	Black	XS	Silk	57.0
2	Under Armour	Dress	Red	M	Wool	127.0
3	NaN	Shoes	Green	M	Cotton	77.0

- Display specific columns

```
In [7]: # Select specific columns (and show the first 3 results)
df[['Brand', 'Category', 'Price']].head(3)
```

```
Out[7]:
```

	Brand	Category	Price
0	New Balance	Dress	182.0
1	New Balance	Jeans	57.0
2	Under Armour	Dress	127.0

- Using loc[] to access data by label

```
In [8]: # Access the row with index label 3
df.loc[3]
```

```
Out[8]: Brand      NaN
Category  Shoes
Color     Green
Size      M
Material  Cotton
Price     77.0
Name: 3, dtype: object
```

```
In [9]: # Select rows 23 to 27 and all columns using label-based indexing
df.loc[23:27]
```

```
Out[9]:
```

	Brand	Category	Color	Size	Material	Price
23	Puma	Jacket	Red	XXL	Wool	158.0
24	Puma	Jacket	Black	L	Denim	125.0
25	Adidas	Jeans	White	XXL	Nylon	89.0
26	Adidas	Dress	Green	L	Polyester	185.0
27	New Balance	Shoes	Blue	M	Cotton	92.0

- Using iloc[] to access data by position

```
In [10]: # Access the row at position 2 (zero-based index)
df.iloc[2]
```

```
Out[10]: Brand      Under Armour
Category      Dress
Color         Red
Size          M
Material      Wool
Price        127.0
Name: 2, dtype: object
```

```
In [11]: # Select rows 1 to 4 and all columns using integer-based indexing
df.iloc[1:5]
```

```
Out[11]:
```

	Brand	Category	Color	Size	Material	Price
1	New Balance	Jeans	Black	XS	Silk	57.0
2	Under Armour	Dress	Red	M	Wool	127.0
3	NaN	Shoes	Green	M	Cotton	77.0
4	NaN	Sweater	White	M	Nylon	113.0

```
In [12]: # Return row 2 and column 5 (In Python indexing starts at 0)
df.iloc[1,4]
```

```
Out[12]: 'Silk'
```

Data Filtering and Selection: boolean indexing

- Equals To

```
In [13]: # Selecting rows where the Brand is "Nike"
df[df['Brand'] == 'Nike'].head()
```

```
Out[13]:
```

	Brand	Category	Color	Size	Material	Price
10	Nike	Jacket	White	XL	Silk	98.0
38	Nike	Dress	Yellow	XS	Silk	137.0
45	Nike	Jacket	Black	XS	Silk	38.0
48	Nike	Jeans	White	XS	Silk	138.0
68	Nike	Shoes	White	L	Cotton	21.0

- Greater Than

```
In [14]: # Selecting rows where the Price is greater than 100
df[df['Price'] > 100].head()
```

Out[14]:

	Brand	Category	Color	Size	Material	Price
0	New Balance	Dress	White	XS	Nylon	182.0
2	Under Armour	Dress	Red	M	Wool	127.0
4	NaN	Sweater	White	M	Nylon	113.0
11	Puma	Jacket	White	XL	Silk	150.0
15	Under Armour	Jacket	Green	XXL	Silk	184.0

- Is In

In [15]: `# Selecting rows where the Category is either "Dress" or "Sweater"`
`df[df['Category'].isin(['Dress', 'Sweater'])].head()`

Out[15]:

	Brand	Category	Color	Size	Material	Price
0	New Balance	Dress	White	XS	Nylon	182.0
2	Under Armour	Dress	Red	M	Wool	127.0
4	NaN	Sweater	White	M	Nylon	113.0
7	Adidas	Dress	Red	XS	Denim	46.0
8	Reebok	Dress	Black	S	Wool	97.0

- Logical AND

In [16]: `# Selecting rows where the Size is "XS" and the Material is "Nylon"`
`df[(df['Size'] == 'XS') & (df['Material'] == 'Nylon')].head()`

Out[16]:

	Brand	Category	Color	Size	Material	Price
0	New Balance	Dress	White	XS	Nylon	182.0
42	Adidas	Jacket	Black	XS	Nylon	173.0
128	Nike	Dress	Red	XS	Nylon	133.0
195	Puma	Sweater	White	XS	Nylon	93.0
215	Under Armour	Dress	Yellow	XS	Nylon	139.0

- Logical OR

In [17]: `# Select rows where Price is less than or equal to 50 or the Color is "Black"`
`df[(df['Price'] <= 50) | (df['Color'] == 'Black')].head()`

Out[17]:

	Brand	Category	Color	Size	Material	Price
1	New Balance	Jeans	Black	XS	Silk	57.0
5	Reebok	Jacket	Red	XL	Nylon	19.0
6	Puma	Jacket	Red	XXL	Polyester	31.0
7	Adidas	Dress	Red	XS	Denim	46.0
8	Reebok	Dress	Black	S	Wool	97.0

Data Manipulation by Adding, Deleting, Renaming columns

- Adding a new column

In this code, a new column named 'Discounted_Price' is created in the DataFrame 'df'. The values in this new column are calculated by multiplying the values in the 'Price' column by 0.5. This operation effectively applies a 50% discount to each price value. After executing this code, if we display the first few rows of the DataFrame using `df.head()`, we'll see the original 'Price' column alongside the newly created 'Discounted_Price' column, which contains the discounted prices.

```
In [18]: df['Discounted_Price'] = df['Price'] * 0.5
df.head()
```

Out[18]:

	Brand	Category	Color	Size	Material	Price	Discounted_Price
0	New Balance	Dress	White	XS	Nylon	182.0	91.0
1	New Balance	Jeans	Black	XS	Silk	57.0	28.5
2	Under Armour	Dress	Red	M	Wool	127.0	63.5
3	NaN	Shoes	Green	M	Cotton	77.0	38.5
4	NaN	Sweater	White	M	Nylon	113.0	56.5

- Deleting an existing column

This code removes the 'Material' column from the DataFrame 'df' using the drop method. The columns parameter specifies the name of the column to be dropped, which is 'Material' in this case. The `inplace=True` argument ensures that the operation is performed on the DataFrame itself and the changes are reflected in place. After executing this code, if we display the first few rows of the DataFrame using `df.head()`, we'll see that the 'Material' column has been removed from the DataFrame.

```
In [19]: df.drop(columns=['Material'], inplace=True)
df.head()
```

Out[19]:

	Brand	Category	Color	Size	Price	Discounted_Price
0	New Balance	Dress	White	XS	182.0	91.0
1	New Balance	Jeans	Black	XS	57.0	28.5
2	Under Armour	Dress	Red	M	127.0	63.5
3	NaN	Shoes	Green	M	77.0	38.5
4	NaN	Sweater	White	M	113.0	56.5

- Renaming an existing column

This code renames the 'Category' column to 'Product_Category' in the DataFrame 'df' using the rename method. The columns parameter specifies a dictionary where the keys are the current column names ('Category') and the values are the new column names ('Product_Category'). The inplace=True argument ensures that the operation is performed on the DataFrame itself and the changes are reflected in place.

In [20]: `df.rename(columns={'Category': 'Product_Category'}, inplace=True)
df.head()`

Out[20]:

	Brand	Product_Category	Color	Size	Price	Discounted_Price
0	New Balance	Dress	White	XS	182.0	91.0
1	New Balance	Jeans	Black	XS	57.0	28.5
2	Under Armour	Dress	Red	M	127.0	63.5
3	NaN	Shoes	Green	M	77.0	38.5
4	NaN	Sweater	White	M	113.0	56.5

Data Cleaning and Preparation

Dealing with Missing Data: dropna(), fillna()

- check for missing values

This method returns the count of missing values for each column in the DataFrame. If a column does not contain any missing values, the sum for that column will be 0. Otherwise, it will return the count of missing values.

In [21]: `df.isnull().sum()`

```
Out[21]: Brand      4
Product_Category  1
Color           0
Size            2
Price           2
Discounted_Price 2
dtype: int64
```

- Drop rows with any missing values

```
In [22]: df.dropna()
```

```
Out[22]:
```

	Brand	Product_Category	Color	Size	Price	Discounted_Price
0	New Balance	Dress	White	XS	182.0	91.0
1	New Balance	Jeans	Black	XS	57.0	28.5
2	Under Armour	Dress	Red	M	127.0	63.5
5	Reebok	Jacket	Red	XL	19.0	9.5
6	Puma	Jacket	Red	XXL	31.0	15.5
...
991	Under Armour	Jacket	Red	XL	18.0	9.0
992	Puma	Sweater	Blue	S	20.0	10.0
993	Puma	Sweater	Green	XXL	15.0	7.5
994	New Balance	Jeans	Black	M	48.0	24.0
996	Puma	Jacket	Red	XXL	110.0	55.0

994 rows × 6 columns

- Fill Missing Values with a Specific Value

```
In [23]: # Fill missing values in the 'Brand' column with 'Unknown'
df.fillna(value={'Brand': 'Unknown'})
```


Out[23]:

	Brand	Product_Category	Color	Size	Price	Discounted_Price
0	New Balance	Dress	White	XS	182.0	91.0
1	New Balance	Jeans	Black	XS	57.0	28.5
2	Under Armour	Dress	Red	M	127.0	63.5
3	Unknown	Shoes	Green	M	77.0	38.5
4	Unknown	Sweater	White	M	113.0	56.5
...
995	Puma	Jeans	Black	NaN	176.0	88.0
996	Puma	Jacket	Red	XXL	110.0	55.0
997	Reebok	Sweater	Blue	NaN	127.0	63.5
998	Unknown	Sweater	Black	XXL	NaN	NaN
999	Unknown	NaN	Yellow	XS	NaN	NaN

1000 rows × 6 columns

- Replacing Missing Values with the Mean

For all instances where the 'Price' data is missing, replace the missing values with the mean price.

```
In [24]: # Calculate the mean value of the 'Price' column
average_price = df['Price'].mean().round(2)

# Fill missing values in the 'Price' column with the mean value (106.26)
df['Price'].fillna(value=average_price, inplace=True)
```

```
In [25]: df
```

Out[25]:

	Brand	Product_Category	Color	Size	Price	Discounted_Price
0	New Balance	Dress	White	XS	182.00	91.0
1	New Balance	Jeans	Black	XS	57.00	28.5
2	Under Armour	Dress	Red	M	127.00	63.5
3	NaN	Shoes	Green	M	77.00	38.5
4	NaN	Sweater	White	M	113.00	56.5
...
995	Puma	Jeans	Black	NaN	176.00	88.0
996	Puma	Jacket	Red	XXL	110.00	55.0
997	Reebok	Sweater	Blue	NaN	127.00	63.5
998	NaN	Sweater	Black	XXL	106.26	NaN
999	NaN	NaN	Yellow	XS	106.26	NaN

1000 rows × 6 columns

Removing Duplicates

- Drop Duplicate Records

Rows are removed that have exactly the same values in all columns as another row in the DataFrame

In [64]: `df.drop_duplicates()`

- Removing duplicates based on selected columns

Duplicates are identified and removed based only on the values in the specified subset of columns.

In [65]: `df.drop_duplicates(subset=['Brand'])`

Convert Data Types

- Convert Price column from Float to Integer

```
In [28]: # print the data type of the 'Price' column before conversion
print("Data type of 'Price' column before conversion:", df['Price'].dtype)
```

Data type of 'Price' column before conversion: float64

```
In [29]: # convert to Integer
df['Price'] = df['Price'].astype(int)
```

```
In [30]: # print the data type of the 'Price' column after conversion
print("Data type of 'Price' column after conversion:", df['Price'].dtype)
```

Data type of 'Price' column after conversion: int32

Data Analysis and Visualization

Descriptive statistics

- `describe()`

Provides summary statistics for numerical columns in a DataFrame

```
In [31]: df.describe().round(2)
```

Out[31]:

	Price	Discounted_Price
count	1000.00	998.00
mean	106.26	53.13
std	53.64	26.85
min	10.00	5.00
25%	59.75	29.62
50%	108.00	54.00
75%	150.00	75.00
max	199.00	99.50

- `mean()`

Calculates the mean of numerical columns

```
In [32]: # Example: Calculate the mean price of products
df['Price'].mean().round(2)
```

Out[32]: 106.26

- `median()`

Calculates the median of numerical columns

```
In [33]: # Example: Calculate the median price of products
df['Price'].median()
```

Out[33]: 108.0

- `std()`

Calculates the standard deviation of numerical column

```
In [34]: # Example: Calculate the standard deviation of prices
df['Price'].std().round(2)
```

```
Out[34]: 53.64
```

- **min() and max()**

Returns the minimum and maximum values of numerical columns

```
In [35]: # Example: Get the minimum price
df['Price'].min()
```

```
Out[35]: 10
```

```
In [36]: # Example: Get the maximum price
df['Price'].max()
```

```
Out[36]: 199
```

- **value_counts()**

Counts the occurrences of each unique value in a column

```
In [37]: df['Product_Category'].value_counts()
```

```
Out[37]: Product_Category
Jacket      190
Shoes       172
Jeans       167
Dress       166
Sweater     160
T-shirt     144
Name: count, dtype: int64
```

- **unique()**

Returns an array of unique values in a column

```
In [38]: df['Color'].unique()
```

```
Out[38]: array(['White', 'Black', 'Red', 'Green', 'Yellow', 'Blue'], dtype=object)
```

- **nunique()**

Returns the number of unique values in a column

```
In [39]: df['Brand'].nunique()
```

```
Out[39]: 6
```

Data Sorting and Grouping

- Sorting Data

```
In [40]: # Sort the DataFrame by the 'Price' column in ascending order.  
df.sort_values(by='Price')
```

```
Out[40]:
```

	Brand	Product_Category	Color	Size	Price	Discounted_Price
519	Puma	T-shirt	Blue	XXL	10	5.0
779	Adidas	Jeans	Red	XXL	10	5.0
115	Under Armour	Jeans	Yellow	XS	10	5.0
271	Reebok	Shoes	Green	XXL	10	5.0
102	Reebok	Jeans	Green	XXL	10	5.0
...
472	New Balance	Shoes	Red	L	198	99.0
498	Adidas	Jacket	Red	XS	198	99.0
194	Puma	Jacket	White	XL	199	99.5
223	Nike	Jeans	White	XS	199	99.5
277	Puma	Dress	Yellow	XS	199	99.5

1000 rows × 6 columns

```
In [41]: # Sort the DataFrame by the 'Price' column in descending order.  
df.sort_values(by='Price', ascending=False)
```

Out[41]:

	Brand	Product_Category	Color	Size	Price	Discounted_Price
223	Nike	Jeans	White	XS	199	99.5
277	Puma	Dress	Yellow	XS	199	99.5
194	Puma	Jacket	White	XL	199	99.5
946	New Balance	Dress	Yellow	S	198	99.0
498	Adidas	Jacket	Red	XS	198	99.0
...
779	Adidas	Jeans	Red	XXL	10	5.0
368	Under Armour	Shoes	Yellow	M	10	5.0
102	Reebok	Jeans	Green	XXL	10	5.0
115	Under Armour	Jeans	Yellow	XS	10	5.0
529	Adidas	Jeans	Black	S	10	5.0

1000 rows × 6 columns

• Grouping Data

In [42]: *# Group DataFrame by 'Brand' column and calculate average price for each brand*
`df.groupby('Brand')['Price'].mean().round(2)`

Out[42]:

Brand	
Adidas	104.00
New Balance	115.10
Nike	102.06
Puma	106.14
Reebok	106.49
Under Armour	104.16

Name: Price, dtype: float64

In [43]: *# Group DataFrame by 'Color' column and count the number of items for each color*
`df.groupby('Color').size()`

Out[43]:

Color	
Black	163
Blue	163
Green	162
Red	168
White	171
Yellow	173

dtype: int64

In [44]: *# Calculate the total sales (sum of prices) for each category*
`df.groupby('Product_Category')['Price'].sum()`

```
Out[44]: Product_Category
Dress      19008
Jacket      19466
Jeans       16825
Shoes       18674
Sweater     17122
T-shirt     15057
Name: Price, dtype: int32
```

```
In [45]: # Calculate the mean price for each combination of brand and category
df.groupby(['Brand', 'Product_Category'])['Price'].mean().round(2)
```

```
Out[45]: Brand      Product_Category
Adidas      Dress      113.37
            Jacket     100.46
            Jeans       92.10
            Shoes     111.96
            Sweater     84.79
            T-shirt    122.67
New Balance Dress     127.67
            Jacket     110.33
            Jeans     114.75
            Shoes     103.35
            Sweater    123.32
            T-shirt    115.46
Nike        Dress      97.68
            Jacket     102.30
            Jeans     106.45
            Shoes     101.25
            Sweater    103.61
            T-shirt    100.96
Puma        Dress     121.33
            Jacket     121.19
            Jeans      97.79
            Shoes     117.78
            Sweater     84.37
            T-shirt     89.54
Reebok      Dress     107.08
            Jacket      93.61
            Jeans     122.54
            Shoes     102.63
            Sweater    116.92
            T-shirt     85.90
Under Armour Dress    120.27
            Jacket     88.91
            Jeans      81.93
            Shoes     113.00
            Sweater    125.22
            T-shirt    109.48
Name: Price, dtype: float64
```

```
In [46]: # Calculate both the mean and sum of prices for each brand
df.groupby('Brand')['Price'].agg(['mean', 'sum']).round(2)
```

Out[46]:

	mean	sum
Brand		
Adidas	104.00	17160
New Balance	115.10	18761
Nike	102.06	16738
Puma	106.14	17831
Reebok	106.49	16826
Under Armour	104.16	18540

Teaser - Data Visualization with Matplotlib

Matplotlib is a Python library used for creating static, interactive, and animated visualizations in Python. It offers a wide range of plotting functions for various types of charts, including line plots, bar charts, histograms, scatter plots, and more. Matplotlib can work seamlessly with Pandas

- Import Matplotlib

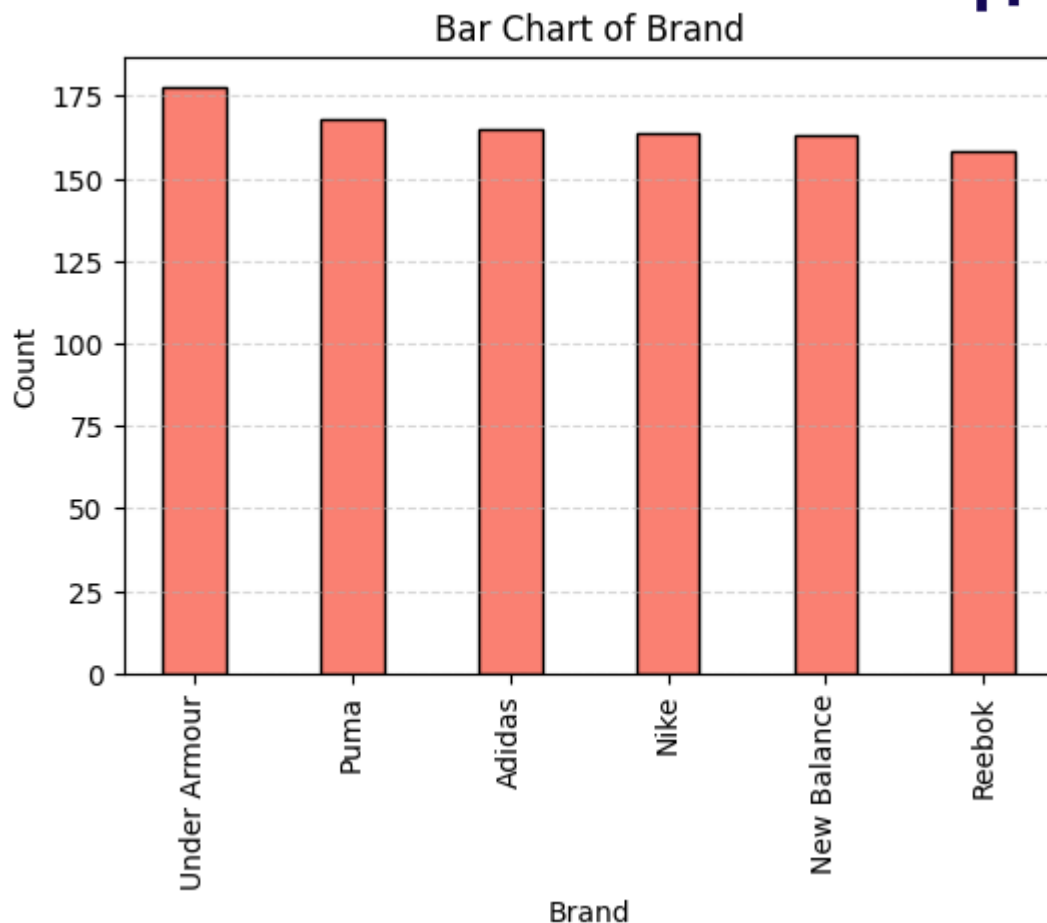
```
In [47]: import matplotlib.pyplot as plt
```

- Bar chart

This code snippet creates a bar chart to display the counts of different brands in a DataFrame. It customizes the appearance of the chart by setting the color, width, and edge color of the bars, as well as adding labels to the axes and a title. Additionally, it includes a dashed grid on the y-axis for better visualization.

```
In [48]: # Count the number of different brands in the DataFrame
brand_counts = df['Brand'].value_counts()

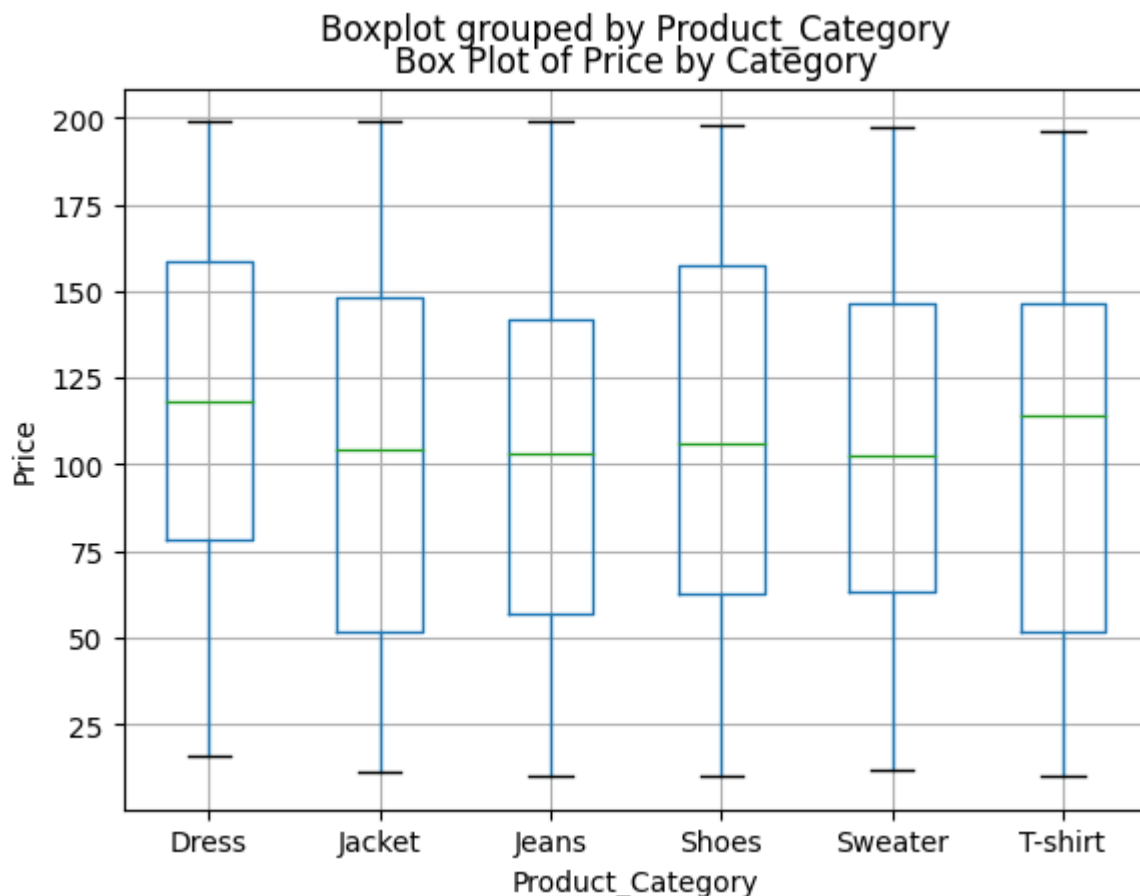
# Create a bar chart with custom settings
plt.figure(figsize=(6, 4)) # Set the figure size
brand_counts.plot(kind='bar', color='salmon', width=0.4, edgecolor='black',
                  linewidth=1)
# Set label for x-axis
plt.xlabel('Brand')
# Set label for y-axis
plt.ylabel('Count')
# Set title of the chart
plt.title('Bar Chart of Brand')
# Add grid lines on the y-axis
plt.grid(axis='y', linestyle='--', alpha=0.5)
# Display the chart
plt.show()
```

- **Box plot**

This code generates a box plot to visualize the distribution of prices across different categories in the DataFrame. The data is grouped by the 'Category' column, and for each category, a box plot of the 'Price' column is displayed. The x-axis represents the categories, while the y-axis represents the prices. The plot is labeled with appropriate axis labels and a title to provide context and interpretation.

```
In [49]: # Create a box plot of Price by Category
df.boxplot(column='Price', by='Product_Category')
# Set labels for x-axis and y-axis
plt.xlabel('Product_Category')
plt.ylabel('Price')
# Set title of the plot
plt.title('Box Plot of Price by Category')
# Display the plot
plt.show()
```



- Line Chart

This code calculates the average prices for different clothing categories in the DataFrame and then creates a line plot to visualize these average prices. Each category is represented on the x-axis, and the corresponding average price is shown on the y-axis. The marker 'o' indicates individual data points, and the line connects these points to illustrate trends. The plot is titled "Average Prices for Clothing Categories" and includes labels for both axes. Finally, grid lines are added to aid in interpretation.

```
In [59]: # Calculate average prices for all clothing items
average_prices = df.groupby('Product_Category')['Price']
average_prices = average_prices.mean().sort_values(ascending=False)

# Plot average prices as a line chart
average_prices.plot(kind='line', marker='o', color='orange')
# Set title and labels for the plot
plt.title('Average Prices for Clothing Categories')
plt.xlabel('Product_Category')
plt.ylabel('Average Price')
# Rotate x-axis labels for better readability
plt.xticks(rotation=45)
# Enable gridlines
plt.grid(True)
# Display the plot
plt.show()
```



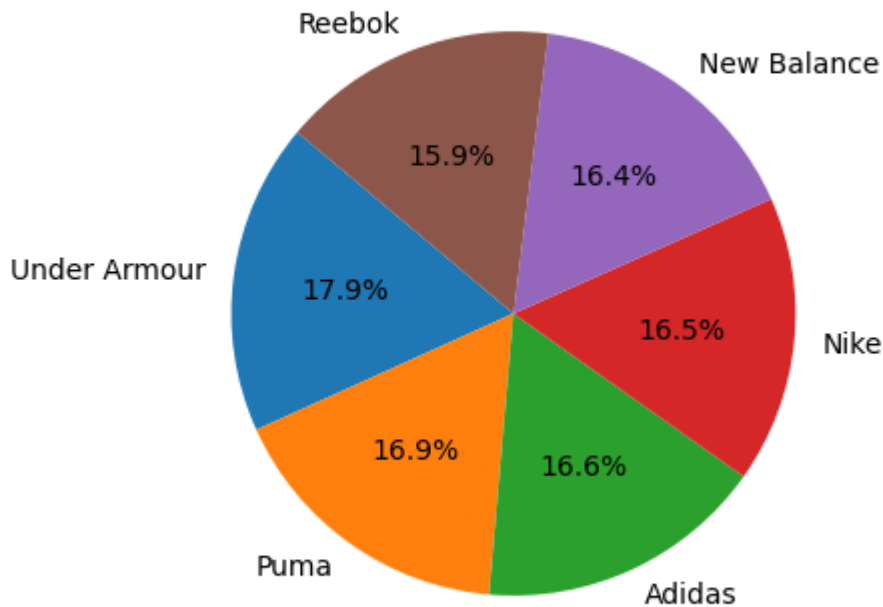
- Pie Chart

This code creates a pie chart that visualizes the distribution of different brands in the DataFrame. It counts the number of occurrences of each brand and then plots the proportions of these counts as slices in the pie chart.

```
In [60]: # Count the number of different brands in the DataFrame
brand_counts = df['Brand'].value_counts()

# Create a Pie Chart
# Set the figure size
plt.figure(figsize=(4, 6))
plt.pie(brand_counts, labels=brand_counts.index, autopct='%1.1f%%',
        startangle=140)
# Add title to the plot
plt.title('Distribution of Brands in the DataFrame')
# Ensure the pie chart is circular
plt.axis('equal')
# Display the plot
plt.show()
```

Distribution of Brands in the DataFrame



Pivot Tables

Pivot tables in Pandas are a way to summarize and analyze data by restructuring it into a more manageable format. They allow you to aggregate and group data based on one or more columns, applying functions to calculate summary statistics or other transformations

- Pivot table to calculate the average price for each Brand-Category pair

```
In [61]: pivot_table_avg_price = df.pivot_table(index='Brand',  
                                                columns='Product_Category',  
                                                values='Price', aggfunc='mean').round(2)  
print(pivot_table_avg_price)
```

Product_Category	Dress	Jacket	Jeans	Shoes	Sweater	T-shirt
Brand						
Adidas	113.37	100.46	92.10	111.96	84.79	122.67
New Balance	127.67	110.33	114.75	103.35	123.32	115.46
Nike	97.68	102.30	106.45	101.25	103.61	100.96
Puma	121.33	121.19	97.79	117.78	84.37	89.54
Reebok	107.08	93.61	122.54	102.63	116.92	85.90
Under Armour	120.27	88.91	81.93	113.00	125.22	109.48

- Pivot table showing the total count of products of each brand in each category

```
In [62]: pivot_table_count = df.pivot_table(index='Brand',  
                                             columns='Product_Category',  
                                             values='Price', aggfunc='count')  
  
print(pivot_table_count)
```

Product_Category	Dress	Jacket	Jeans	Shoes	Sweater	T-shirt
Brand						
Adidas	30	28	39	25	19	24
New Balance	27	33	24	31	22	26
Nike	22	30	22	32	31	27
Puma	21	32	28	36	27	24
Reebok	36	23	24	19	36	20
Under Armour	30	44	30	28	23	23

- Count of Items by Category and Size

```
In [63]: pivot_table_item_count = df.pivot_table(index='Product_Category',  
                                                  columns='Size', values='Price',  
                                                  aggfunc='count')  
  
pivot_table_item_count
```

Out[63]:

	Size	L	M	S	XL	XS	XXL
--	------	---	---	---	----	----	-----

Product_Category							
	Dress	27	21	27	25	41	25
	Jacket	27	25	38	35	31	34
	Jeans	24	24	20	24	37	37
	Shoes	21	32	31	26	36	26
	Sweater	26	27	26	25	26	29
	T-shirt	15	28	24	32	23	22